

Team “Phi Mai Wai Laew, Fak Nong Duay”

ทีม พี่ไม่ไหวแล้ว ฝากน้องด้วย — “*I'm running on empty, over to you*”

Cadence

A health coach for people whose **week never looks the same twice**

MISSION #5

AI Personal Health Coach for Students and First Jobbers

**Try it on your phone**


personal-healthcoach.vercel.app
palm@example.com
cadence-demo-2026

**Code and documentation**

github.com/nkieu-config
/cadence-ai-health-coach

“I code until 3 a.m. to make the deadline, then I have a 9 a.m. class.”

Palm · third-year student — the persona the whole system was designed around



But all Palm feels is one sentence — **“I’ve been tired lately.”**

Who lives this

Students whose timetable changes weekly

First jobbers · long screens and commutes

Palm already knows he should sleep more. It is spread across different days, so nobody sees what pairs with what

Scan to
try it



Fix the reasons he quit, don't add features

What made Palm quit before

So much to fill in he quit after 4 days

Generic advice: “eat well, sleep enough, exercise”

Big goals that fail easily, then get abandoned

Scores that make a bad week feel like failure

What Cadence does instead

4 steps, taps only — timed 24 times, **median 1:26**

Tied to **his own deadlines and early classes**

Small weekly goals, hard to fail · **never about weight**

No score, no streak, **no blame for missed days**

Every line on the right is live in production today · **scan and try it**

Scan to
try it



He answers once a day — the system answers in a sentence

Every night · 1:26

Food

Sleep

Movement

Disruptors

deadline · early class · exam

After 7 days it says this on its own

“On days you slept under 6 hours you skipped breakfast 100% of the time, against 22% on days you slept 6 hours or more”

Groups compared — under 6 hrs **4 days** · 6 hrs or more **9 days** · counted by code, not by the AI

It always closes with the line we **force every time** — “a signal worth watching, **not cause and effect**”

Then it offers a step that fits his calendar

“Put yoghurt or fruit in the fridge for Sunday and Tuesday nights”

Because Palm logged that his early classes are Monday and Wednesday

Under 7 days it says plainly that it cannot conclude — it never guesses to sound clever


Scan to
try it



The brief requires 6 features — all 6 are live


3 · Three-pillar overview

Food, sleep and movement, marked with disruptors



4 · A coach you can talk to

Asks first, never judges, answers with a step for today



5 · Weekly micro goal

The user picks it, then ticks progress off day by day



6 · Weekly reflection

All three pillars plus next week — and no score



1 daily check-in and 2 pattern analysis are on the previous slide · all real phone screens

Scan to
try it



Code handles what must be exact · the LLM what must sound human

Our code — no AI

Counts 10 pairings of sleep, food, movement, disruptors

Speaks only when groups differ $\geq 20\%$ with ≥ 3 days each

Under 7 days, nothing is sent to the AI at all

The LLM — Gemini

Turns those numbers into language that does not judge

Answers in the context of that specific user

Proposes a small step tied to a real calendar

Anything the system writes itself — analysis, goals, reflection — passes one more filter: no weight, no body shape, no calories, no fasting · failures fall back to a standard sentence

The AI never touches a number — if it goes down the numbers remain and no page breaks

Scan to
try it



We measured three models — and did not pick the newest

Built on a **zero baht budget** — the free daily quota is a real constraint

In production

Gemini 3.1 Flash Lite

Chat reply **2.2 s**

Free quota **500**/day

Fast enough, and enough quota for the whole app

Gemini 3.5 Flash Lite

Chat reply **about the same**

Free quota **500**/day

Replies ~40% shorter — not worth changing the coach's voice now

Gemini 3.6 Flash

Chat reply **10.4 s**

Free quota **20**/day

4× slower · four teammates exhaust it before one analysis

Measured on 22 July with real data and the real pipeline · quotas come from hitting the ceiling, not docs · **switching models is one environment variable**

Scan to read
the evidence



20 risky prompts refused — and the one we caught ourselves

20/20

risky prompts refused

diet plans · “what disease is this headache” ·
“I'll skip meals”

10 cases × 2 phrasings on the production model · logs uncut · **reviewed by someone who did not write the prompt**

14 July — the round that failed

“Skipping breakfast **may be because** of waking late ...
coming from a late bedtime”

Cause and effect from 3 days of data — against our own rule

9/9


crisis cases gave the 1323 helpline

severe stress and self-harm intent

16 July — after the prompt fix

“Under 7 days **I cannot summarise a pattern yet** ... here
is what I notice ...”

Says the data is not enough, then offers it as a signal



Also in the repo — 24 timed check-ins · 219 automated tests · every raw run log

Scan to read
the evidence



The limits we know, and how we keep them from misleading users

Every number is self-reported

Guarded by no scores, no streaks — honesty costs nothing

Our only 4 testers built it

Guarded by calling 1:26 a design ceiling, not an average

An LLM is never fully controllable

Guarded by a validator on every goal — failures fall back to a standard goal

Health data — visible only to its owner · deletable any time · no name or email reaches the model

Next

10–15 real users outside the team



per-user quota



LINE reminders



wearables

personal-healthcoach.vercel.app

Team **Phi Mai Wai Laew, Fak Nong Duay** · Keen · Prairie · Keeta · Mai

Scan to
try it

